

# Real-Time Intelligence: Accelerating Insights with RisingWave

Architectural Strategy & E-Commerce  
Conversion Demo



Savas Gioldasis | Principal Data Engineer

# Moving from Retrospective to Reactive

Traditional batch analytics are too slow for modern conversion optimization. We are validating a stack that delivers immediate insights without the Big Data overhead.

## The Solution



RisingWave – A cloud-native streaming database compatible with PostgreSQL.

## The Advantage



High throughput with significantly lower TCO compared to legacy stacks like Flink.

## The Proof



A live E-Commerce Funnel demo showing real-time View-to-Cart and Cart-to-Buy conversion rates.

# RisingWave: Streaming Made as Simple as SQL

## Traditional Database



Static Data. Batch Queries.

- **Definition:** A cloud-native streaming database that uses SQL as the primary interface.

## RisingWave



Dynamic Streams. Continuous Queries.

- **The Paradigm:** Treats data streams as tables. No complex Java/Scala; just standard SQL queries that update automatically.

- **Capabilities:** Powering real-time dashboards, ad-hoc queries, and alert systems immediately upon ingestion.

# Strategic Comparison: RisingWave vs. Apache Flink

## Apache Flink (The Incumbent)



- **Pros:** Mature ecosystem, massive scale.
- **Cons:** High operational complexity, heavy infrastructure footprint, requires specialized Java/Scala talent.

## RisingWave (The Agile Challenger)



- **Pros:** PostgreSQL Compatible (leverages existing SQL talent), Cloud-native architecture (separation of compute/storage), significantly lower TCO.
- **Cons:** Newer ecosystem.

**Verdict:** RisingWave offers faster time-to-market for agile real-time applications.

# The Business Case: E-Commerce Conversion Funnel

## Viewers

Unique page views  
(src\_page.sql)

## Carters

Users adding items to cart  
(src\_cart.sql)

## Purchasers

Completed transactions  
(src\_purchase.sql)

**The Goal:**  
Instantly calculate  
`view_to_cart_rate` and  
`cart_to_buy_rate` to  
identify drop-offs in  
real-time.

# Live Visualization: The Modern React Dashboard

Viewer Count:  
**12,450**

Cart Rate:  
**45%**

Purchase Rate:  
**12%**

## Producer Controls

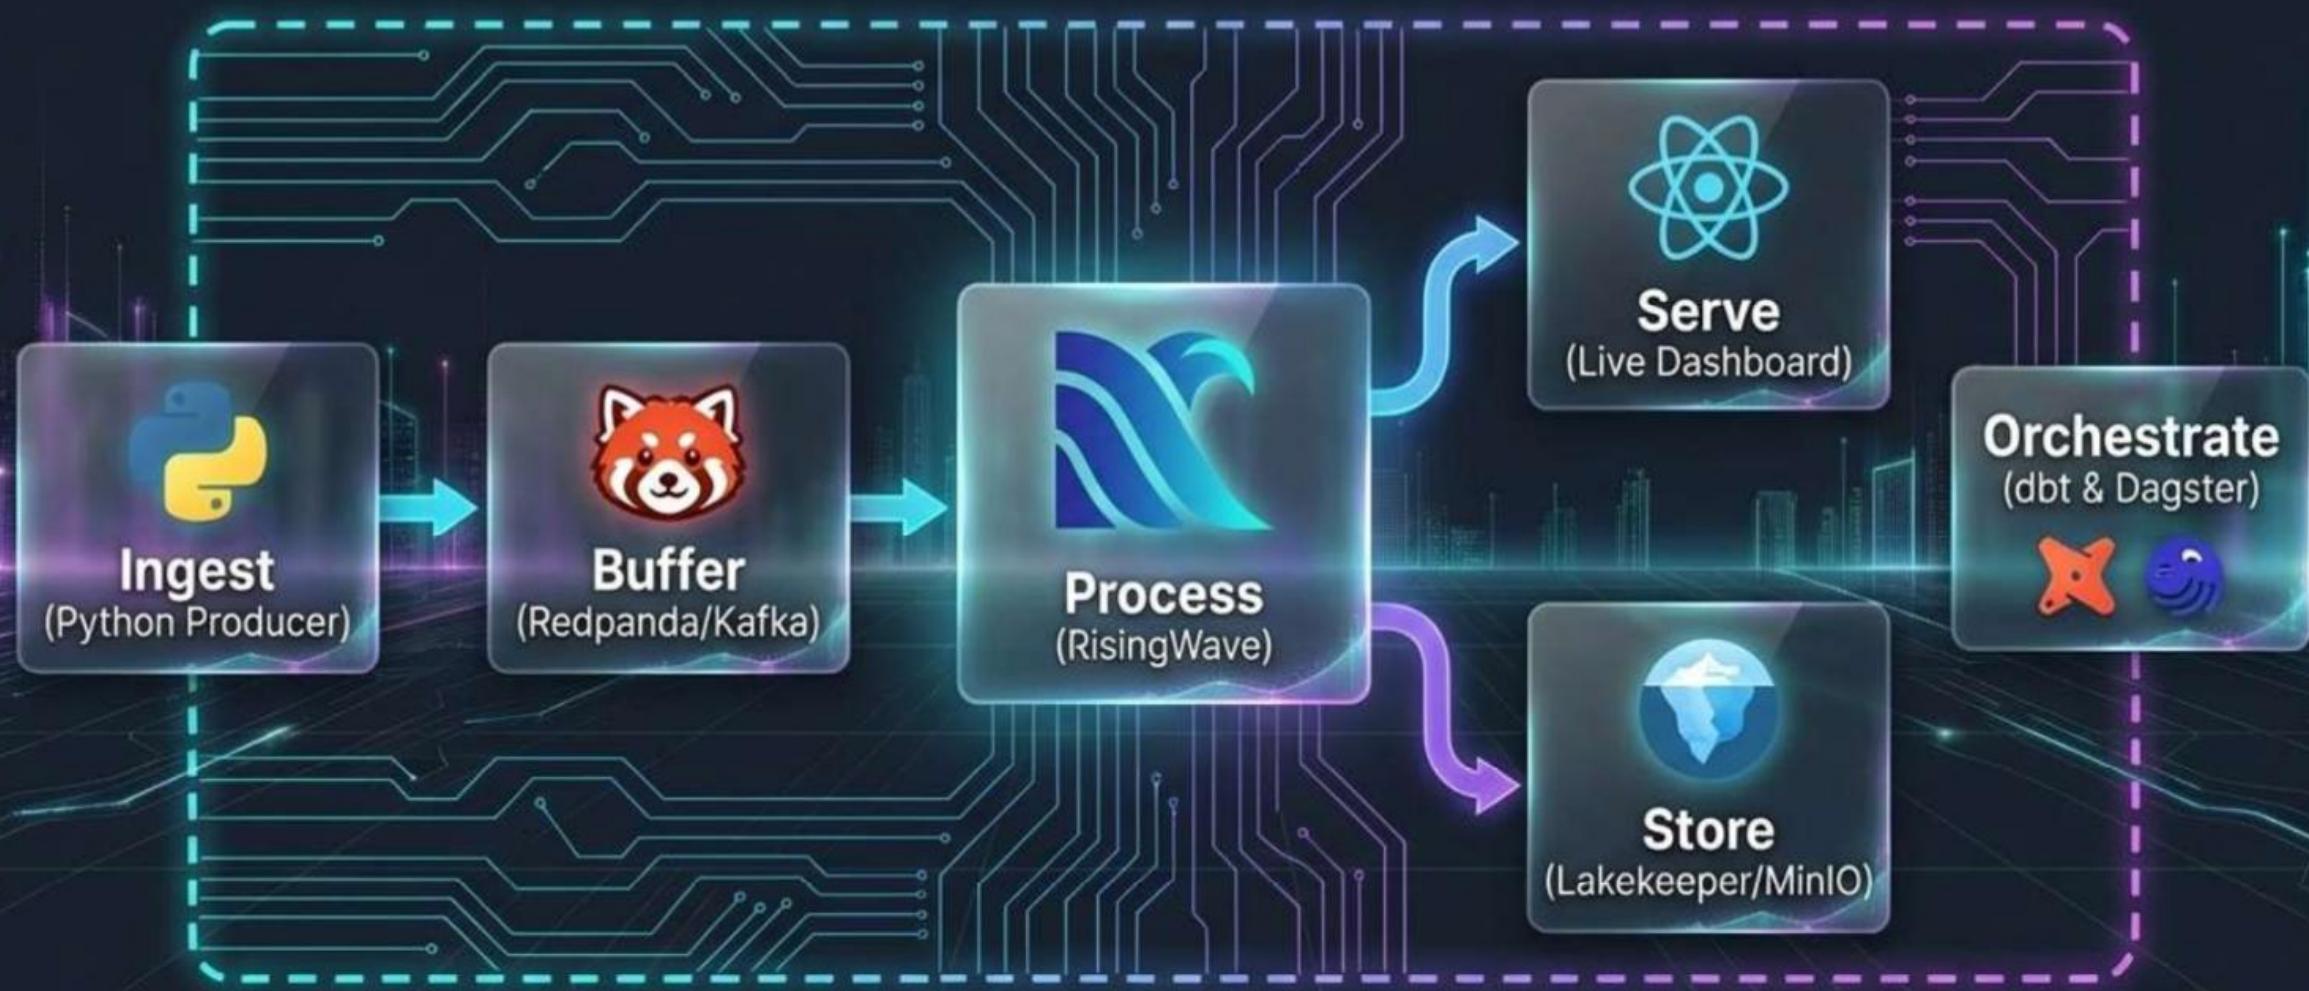
▶ Start

■ Stop

**Real-time Updates:**  
Millisecond-level latency.

**Tech Stack:**  
React Frontend + FastAPI  
Backend + RisingWave.

# High-Level Architecture



# The Data Journey: From Event to Materialized View



# Orchestration & Quality Assurance

## Orchestration (Dagster)



Manages pipeline lifecycle  
and ensures reliability.

## Transformation (dbt)

```
SELECT *  
FROM {{ source('risingwave', 'page_views') }}
```

Applies version control and  
testing to streaming logic.

# The Persistence Layer: Apache Iceberg & Lakekeeper



# Advanced Analytics: DuckDB & Spark Integration



DuckDB: Ad-hoc queries on Iceberg tables  
(scripts/query\_raw\_iceberg.py)

Spark: Complex user behavior modeling and  
historical research.

# Key Strategic Takeaways



**Simplicity:** Complex streaming logic implemented using standard SQL and dbt.



**Performance:** Achieved millisecond latency using Redpanda and RisingWave's compute-storage separation.



**Scalability:** Production-ready architecture separating speed (RisingWave) from storage (Iceberg/MinIO).

# Ready to Scale Real-Time Intelligence

The architecture demonstrated is modular, scalable,  
and ready for pilot deployment.



**Savas Gioldasis**  
Principal Data Engineer